

## ASL\_detection\_project

This document contains a complete, runnable ASL detection project you can download and run locally. It assumes your dataset follows this structure:

data/

train/

A/...

B/...

... (all 29 classes)

test/

A/...

B/...

...

---

### Project structure

ASL\_detection\_project/

├─ data/           # (your dataset - not included)

├─ models/

├─ src/

|   ├─ train.py

|   ├─ evaluate.py

|   └─ predict.py

├─ requirements.txt

└─ README.md

---

### requirements.txt

tensorflow>=2.10

numpy

matplotlib

scikit-learn

opencv-python

pillow

---

## README.md

### # ASL Detection

#### ## Setup

1. Create a Python 3.8+ virtual environment and install requirements:

```
``bash
```

```
python -m venv venv
```

```
source venv/bin/activate # or .\venv\Scripts\activate on Windows
```

```
pip install -r requirements.txt
```

2. Put your dataset in data/train and data/test with one subfolder per class (29 classes).

3. Train the model:

```
python src/train.py --data_dir data --img_size 128 --batch_size 32 --epochs 30
```

4. Evaluate:

```
python src/evaluate.py --data_dir data --img_size 128 --batch_size 32 --model_path  
models/asl_mobilenetv2.h5
```

5. (Optional) Run webcam demo:

```
python src/predict.py --model_path models/asl_mobilenetv2.h5 --img_size 128
```

```
---
```

#### ## src/train.py

```
``python
```

```
"""Train a transfer-learning ASL classifier using MobileNetV2.
```

```
Saves model to models/asl_mobilenetv2.h5 and a saved_model folder.
```

```
"""
```

```
import argparse
```

```
import os
```

```
from pathlib import Path

import tensorflow as tf

from tensorflow.keras import layers, models

from tensorflow.keras.applications import MobileNetV2

from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
```

```
def get_datasets(data_dir, img_size, batch_size):
```

```
    train_dir = os.path.join(data_dir, 'train')
```

```
    val_dir = os.path.join(data_dir, 'test')
```

```
    train_ds = tf.keras.preprocessing.image_dataset_from_directory(
        train_dir,
        labels='inferred',
        label_mode='categorical',
        batch_size=batch_size,
        image_size=(img_size, img_size),
        shuffle=True
    )
```

```
    val_ds = tf.keras.preprocessing.image_dataset_from_directory(
        val_dir,
        labels='inferred',
        label_mode='categorical',
        batch_size=batch_size,
        image_size=(img_size, img_size),
        shuffle=False
    )
```

```
    class_names = train_ds.class_names
```

```
# Prefetch for performance

AUTOTUNE = tf.data.AUTOTUNE

train_ds = train_ds.prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.prefetch(buffer_size=AUTOTUNE)
```

```
return train_ds, val_ds, class_names
```

```
def build_model(num_classes, img_size, fine_tune_at=100):
```

```
    base_model = MobileNetV2(input_shape=(img_size, img_size, 3), include_top=False,
weights='imagenet')
```

```
    base_model.trainable = True
```

```
    # Freeze early layers
```

```
    for layer in base_model.layers[:fine_tune_at]:
```

```
        layer.trainable = False
```

```
    inputs = layers.Input(shape=(img_size, img_size, 3))
```

```
    x = layers.Rescaling(1./127.5, offset=-1)(inputs) # MobileNetV2 preprocessing
```

```
    # Data augmentation (simple)
```

```
    data_augmentation = tf.keras.Sequential([
```

```
        layers.RandomFlip('horizontal'),
```

```
        layers.RandomRotation(0.08),
```

```
        layers.RandomZoom(0.08),
```

```
    ])
```

```
    x = data_augmentation(x)
```

```
    x = base_model(x, training=False)
```

```
    x = layers.GlobalAveragePooling2D()(x)
```

```
    x = layers.Dropout(0.3)(x)
```

```
outputs = layers.Dense(num_classes, activation='softmax')(x)
```

```
model = models.Model(inputs, outputs)
```

```
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4),  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])  
return model
```

```
def main(args):
```

```
    os.makedirs('models', exist_ok=True)
```

```
    train_ds, val_ds, class_names = get_datasets(args.data_dir, args.img_size, args.batch_size)  
    num_classes = len(class_names)
```

```
    model = build_model(num_classes, args.img_size, fine_tune_at=100)  
    model.summary()
```

```
    checkpoint = ModelCheckpoint('models/asl_mobilenetv2.h5', monitor='val_accuracy',  
                                save_best_only=True, verbose=1)
```

```
    early = EarlyStopping(monitor='val_loss', patience=6, restore_best_weights=True)
```

```
    reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr=1e-7)
```

```
    history = model.fit(  
        train_ds,  
        validation_data=val_ds,  
        epochs=args.epochs,  
        callbacks=[checkpoint, early, reduce_lr]  
    )
```

```

# Save final model (also saved_model format)

model.save('models/asl_mobilenetv2.h5')

model.save('models/asl_mobilenetv2_saved')


if __name__ == '__main__':
    parser = argparse.ArgumentParser()
    parser.add_argument('--data_dir', type=str, default='data')
    parser.add_argument('--img_size', type=int, default=128)
    parser.add_argument('--batch_size', type=int, default=32)
    parser.add_argument('--epochs', type=int, default=30)
    args = parser.parse_args()
    main(args)

```

---

### **src/evaluate.py**

```

"""Evaluate model on test set and show confusion matrix + per-class accuracy."""

import argparse
import os
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report
import itertools
import tensorflow as tf

def get_dataset(data_dir, img_size, batch_size):
    test_dir = os.path.join(data_dir, 'test')
    ds = tf.keras.preprocessing.image_dataset_from_directory(
        test_dir,
        labels='inferred',
        label_mode='categorical',

```

```

        batch_size=batch_size,
        image_size=(img_size, img_size),
        shuffle=False
    )
    return ds

```

```

def plot_confusion_matrix(cm, class_names, out_path='confusion_matrix.png'):

```

```

    fig = plt.figure(figsize=(12, 10))
    plt.imshow(cm, interpolation='nearest')
    plt.title('Confusion matrix')
    plt.colorbar()
    tick_marks = np.arange(len(class_names))
    plt.xticks(tick_marks, class_names, rotation=90)
    plt.yticks(tick_marks, class_names)

    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], 'd'),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

```

```

    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.tight_layout()
    fig.savefig(out_path)
    print(f"Saved confusion matrix to {out_path}")

```

```

def main(args):

```

```

    ds = get_dataset(args.data_dir, args.img_size, args.batch_size)

```

```

class_names = ds.class_names

model = tf.keras.models.load_model(args.model_path)

y_true = []
y_pred = []

for images, labels in ds:
    preds = model.predict(images)
    y_true.extend(np.argmax(labels.numpy(), axis=1).tolist())
    y_pred.extend(np.argmax(preds, axis=1).tolist())

cm = confusion_matrix(y_true, y_pred)
plot_confusion_matrix(cm, class_names, out_path='models/confusion_matrix.png')

print('\nClassification report:\n')
print(classification_report(y_true, y_pred, target_names=class_names))

# Per-class accuracy
class_acc = cm.diagonal() / cm.sum(axis=1)
for name, acc in zip(class_names, class_acc):
    print(f'{name}: {acc:.3f}')

if __name__ == '__main__':
    parser = argparse.ArgumentParser()
    parser.add_argument('--data_dir', type=str, default='data')
    parser.add_argument('--img_size', type=int, default=128)
    parser.add_argument('--batch_size', type=int, default=32)
    parser.add_argument('--model_path', type=str, default='models/asl_mobilenetv2.h5')
    args = parser.parse_args()

```



```
main(args)
```

---

### **src/predict.py**

```
"""Simple webcam demo that predicts ASL signs in real-time.
```

```
Press 'q' to quit.
```

```
"""
```

```
import argparse
```

```
import cv2
```

```
import numpy as np
```

```
import tensorflow as tf
```

```
def main(args):
```

```
    model = tf.keras.models.load_model(args.model_path)
```

```
    img_size = args.img_size
```

```
    # We need class names. If the model was trained with Keras Dataset, we will read from data/train
```

```
    import os
```

```
    classes = sorted(os.listdir(os.path.join(args.data_dir, 'train')))
```

```
    cap = cv2.VideoCapture(0)
```

```
    if not cap.isOpened():
```

```
        print('Cannot open webcam')
```

```
        return
```

```
    while True:
```

```
        ret, frame = cap.read()
```

```
        if not ret:
```

```
            break
```

```

# Flip and resize
frame_rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
img = cv2.resize(frame_rgb, (img_size, img_size))
img = img.astype('float32') / 127.5 - 1.0 # MobileNetV2 preprocessing
input_arr = np.expand_dims(img, axis=0)

preds = model.predict(input_arr)
idx = np.argmax(preds[0])
label = classes[idx]
prob = preds[0][idx]

# Overlay
cv2.putText(frame, f'{{label}} ({{prob:.2f}})', (10, 30), cv2.FONT_HERSHEY_SIMPLEX, 1, (0,255,0), 2)
cv2.imshow('ASL demo', frame)

if cv2.waitKey(1) & 0xFF == ord('q'):
    break

cap.release()
cv2.destroyAllWindows()

if __name__ == '__main__':
    parser = argparse.ArgumentParser()
    parser.add_argument('--model_path', type=str, default='models/asl_mobilenetv2.h5')
    parser.add_argument('--img_size', type=int, default=128)
    parser.add_argument('--data_dir', type=str, default='data')
    args = parser.parse_args()
    main(args)

```

---

## Notes & Tips

- If your dataset is small, consider heavier augmentation or K-fold cross validation.
  - To reduce model size for deployment, try `tf.keras.models.save_model(..., include_optimizer=False)` and/or convert to TFLite.
  - If training on CPU is slow, use a GPU runtime (local GPU or Colab).
- 

If you'd like, I can also:

- Convert this into a downloadable ZIP.
- Add a Colab notebook that mounts your Google Drive dataset and trains on GPU.
- Replace MobileNetV2 with a custom CNN if you prefer a smaller-from-scratch model.

Tell me which of the above you'd like next.